

GENERATED PRODUCT IMAGES

1,600 images across 8 product categories

SHARED EVALUATION SUBSET

200 images: first 25 indexed per category

HUMAN JUDGMENTS

30 participants × 200 images

6,000 participant responses

Five meaning-label shares per image

Plurality labels retain exact ties

PROMPTED MODEL LABELS

9 vision-language models

One selected label per image

Qwen image / category controls

Prompt-order and repeat-run checks

SUPERVISED COMPARISONS

Category-only and global priors

Soft-target / plurality-label LoRA

Fit using training responses only

Balanced and category-held-out folds

SELECTED-LABEL AGREEMENT

Tie-aware accuracy and human vote-share support

Paired, category-stratified image-bootstrap intervals

PRESERVED PROBABILITY DISTRIBUTIONS

Adapter / prior probabilities vs. human response shares

Summed Brier score · soft cross-entropy · JSD